

Support Vector Machines (SVMs)

Support Vector Machines (SVMs) are a type of supervised learning algorithm used for classification and regression analysis. SVMs can be tuned with hyperparameters to optimize their performance on a given dataset. Here are some of the most common hyperparameters for SVMs and their meanings:

1. **C:** The C hyperparameter controls the trade-off between maximizing the margin and minimizing the classification error.

- When C is small, the SVM tries to maximize the margin between the decision boundary and the training examples, even if that means misclassifying some of the training examples. This can lead to a more robust model that generalizes well to new, unseen data. However, if C is too small, the SVM may be underfitting the training data, resulting in poor performance on the training and validation sets.
- When C is large, the SVM tries to correctly classify as many training examples as possible, even if that means having a smaller margin. This can lead to a more complex decision boundary that can overfit the training data, resulting in poor performance on the validation set. However, if C is too large, the SVM may be overfitting the training data, resulting in poor performance on the validation set.

2. **kernel:** The kernel hyperparameter specifies the kernel function to be used for mapping the input features into a higher-dimensional space. Some of the popular kernel functions used in SVMs are linear, polynomial, and radial basis function (RBF).

- The kernel function is a similarity measure between two data points, and it determines the shape of the decision boundary that separates the classes in the transformed feature space. SVMs use the kernel trick to avoid explicitly computing the coordinates of the data points in the high-dimensional space, which can be computationally expensive or even impossible if the number of dimensions is very large.
- There are several types of kernel functions that can be used in SVMs. Here are some of the most commonly used kernel functions and their characteristics:
- **Linear kernel:** The linear kernel is the simplest kernel function, and it is used when the data is linearly separable. The decision boundary is a straight line that separates the two classes.
- **Polynomial kernel:** The polynomial kernel is used when the data is not linearly separable. The degree of the polynomial kernel determines the complexity of the decision boundary, and a higher degree can result in overfitting.

- **Radial basis function (RBF) kernel:** The RBF kernel is a popular kernel function that is used when the data is not linearly separable. The RBF kernel has a smooth, non-linear decision boundary that can fit complex patterns in the data. The RBF kernel has a parameter γ which controls the width of the kernel and determines the flexibility of the decision boundary. A smaller value of γ results in a smoother decision boundary, while a larger value of γ results in a more complex decision boundary that can overfit the data.
- **Sigmoid kernel:** The sigmoid kernel is another kernel function that is used when the data is not linearly separable. The sigmoid kernel has a similar shape to the logistic function, and it is suitable for binary classification problems. However, the sigmoid kernel is sensitive to the choice of hyperparameters and can lead to poor performance if not properly tuned.
- **Choosing the right kernel function** depends on the nature of the data and the problem at hand. It's a good idea to try different kernel functions and compare their performance on a validation set to determine the best kernel function for a given problem.

3. **γ :** The γ hyperparameter controls the shape of the decision boundary. A smaller value of γ results in a smoother decision boundary, while a larger value of γ results in a more complex decision boundary that is more likely to overfit.

4. **degree:** The degree hyperparameter specifies the degree of the polynomial kernel function, if the kernel used is polynomial. This hyperparameter is usually set to 3 by default.

5. **class_weight:** The class_weight hyperparameter can be used to handle imbalanced classes in the dataset. It assigns a weight to each class, with larger weights given to minority classes and smaller weights given to majority classes.

To tune these hyperparameters, you can use a technique called **grid search**, which involves training the SVM model with different combinations of hyperparameter values and evaluating their performance on a validation set. You can then choose the combination of hyperparameters that gives the best performance on the validation set.

Logistic regression

Logistic regression is a binary classification algorithm that predicts the probability of an input belonging to a particular class (usually 0 or 1). The hyperparameters in logistic regression are tuning parameters that determine the performance of the model. Here are some of the most important hyperparameters in logistic regression:

Penalty (C): The penalty parameter C controls the regularization strength of the model. A higher value of C will reduce the regularization strength, which can lead to overfitting. On the

other hand, a lower value of C will increase the regularization strength, which can lead to underfitting. Typically, a value of C between 0.1 and 100 is used.

Solver: The solver parameter determines the algorithm used to optimize the logistic regression objective function. Some common solvers include 'lbfgs', 'newton-cg', 'sag', and 'liblinear'. The choice of solver will depend on the size of the dataset and the complexity of the problem.

Class weight: The class weight parameter is used to balance the weight of the classes in the dataset. If the dataset is imbalanced, meaning that one class has many more samples than the other, then setting the class weight parameter can help improve the performance of the model.

Max iterations: The max_iter parameter determines the maximum number of iterations allowed for the solver to converge. If the model does not converge within the specified number of iterations, it will return an error.

Regularization: Logistic regression can use different regularization methods, such as L1 or L2 regularization, to avoid overfitting. The regularization parameter determines the strength of the regularization.

Dual: The dual parameter determines whether to solve the dual or primal optimization problem. The dual problem is used when there are more features than samples, while the primal problem is used when there are more samples than features.

Gradient boosting

Gradient boosting regression is a machine learning algorithm that uses an ensemble of decision trees to predict a continuous target variable. Here are some of the most important hyperparameters in Gradient Boosting Regression:

n_estimators: This hyperparameter controls the number of trees in the ensemble. Increasing the number of trees can improve the performance of the model, but can also increase the risk of overfitting.

learning_rate: The learning_rate hyperparameter controls the rate at which the model adapts to the data. A smaller learning rate results in a slower learning process, but can lead to better generalization. Conversely, a higher learning rate will result in a faster learning process, but can lead to overfitting.

max_depth: This hyperparameter controls the maximum depth of the individual decision trees in the ensemble. Increasing the maximum depth can improve the performance of the model, but can also increase the risk of overfitting.

min_samples_split: The min_samples_split hyperparameter controls the minimum number of samples required to split an internal node. Increasing this value can lead to simpler models and reduce overfitting.

loss function: The loss function hyperparameter determines the objective function that the algorithm optimizes. Common loss functions include 'ls' for least squares regression and 'lad' for least absolute deviation regression.

subsample: This hyperparameter controls the fraction of the training data that is used to train each individual tree in the ensemble. A value less than 1.0 can reduce the risk of overfitting by introducing randomness into the model.

max_features: The max_features hyperparameter controls the maximum number of features used to split each node in the decision trees. A smaller value can reduce the risk of overfitting and improve generalization.

Random Forest

Random Forest is an ensemble machine learning algorithm that combines multiple decision trees to make predictions. Here are some of the most important hyperparameters in Random Forest:

n_estimators: This hyperparameter controls the number of decision trees in the ensemble. Increasing the number of trees can improve the performance of the model, but can also increase the risk of overfitting.

max_depth: This hyperparameter controls the maximum depth of each decision tree in the ensemble. Increasing the maximum depth can improve the performance of the model, but can also increase the risk of overfitting.

min_samples_split: The min_samples_split hyperparameter controls the minimum number of samples required to split an internal node. Increasing this value can lead to simpler models and reduce overfitting.

min_samples_leaf: The min_samples_leaf hyperparameter controls the minimum number of samples required to be at a leaf node. Increasing this value can lead to simpler models and reduce overfitting.

max_features: The max_features hyperparameter controls the maximum number of features used to split each node in the decision trees. A smaller value can reduce the risk of overfitting and improve generalization.

bootstrap: The bootstrap hyperparameter controls whether or not to use bootstrapping when sampling the training data. Bootstrapping can introduce randomness into the model and reduce the risk of overfitting.

criterion: The criterion hyperparameter controls the quality of the split. The two most common criteria are 'gini' and 'entropy'.

Price prediction

*The main difference between **normal price prediction**, **time series price prediction**, and **geographical changes price prediction** is the type of data being used and the modeling techniques that are applied.*

Normal price prediction: Normal price prediction refers to predicting the price of a product or service based on a set of features that are relevant to the price. For example, features such as product quality, brand reputation, consumer demand, and competition can be used to predict the price of a product. Normal price prediction is often done using machine learning algorithms such as regression, random forest, or support vector regression.

Time series price prediction: Time series price prediction refers to predicting the price of a product or service over time. Time series data is typically collected at regular intervals and includes historical prices, along with other relevant time-varying features such as seasonality, trends, and cyclical patterns. Time series price prediction is often done using statistical techniques such as autoregressive integrated moving average (ARIMA), seasonal ARIMA, or exponential smoothing.

Geographical changes price prediction: Geographical changes price prediction refers to predicting the price of a product or service in different regions or locations. Geographical price prediction involves identifying the factors that influence the price of a product or service in a specific region, such as local demand, supply, competition, and regulatory environment. Geographical price prediction is often done using machine learning algorithms such as regression, random forest, or support vector regression, with additional geographical data such as location, population density, and economic indicators included as features.

the main difference between these three types of price prediction is the type of data being used and the modeling techniques that are applied.

- Normal price prediction uses features that are relevant to the price,
- time series price prediction uses historical prices and time-varying features,
- geographical changes price prediction uses regional data and features.